

S1364040 Assignment-02(updated)

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 cd '/content/drive/MyDrive/Colab Notebooks/'
```

/content/drive/MyDrive/Colab Notebooks

1. Exploratory Data Analysis (EDA)

```
1 # Import necessary libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from tensorflow.keras.utils import to_categorical
7 import random
8
```

```
1 # Load the image data and labels
2 images = np.load('images.npy')
3 labels = pd.read_csv('Labels.csv')
4
5 # Check the structure of the loaded data
6 print("Images shape:", images.shape) # Check dimensions of the images
7 print("Labels head:\n", labels.head()) # Display the first few rows of labels
8
```

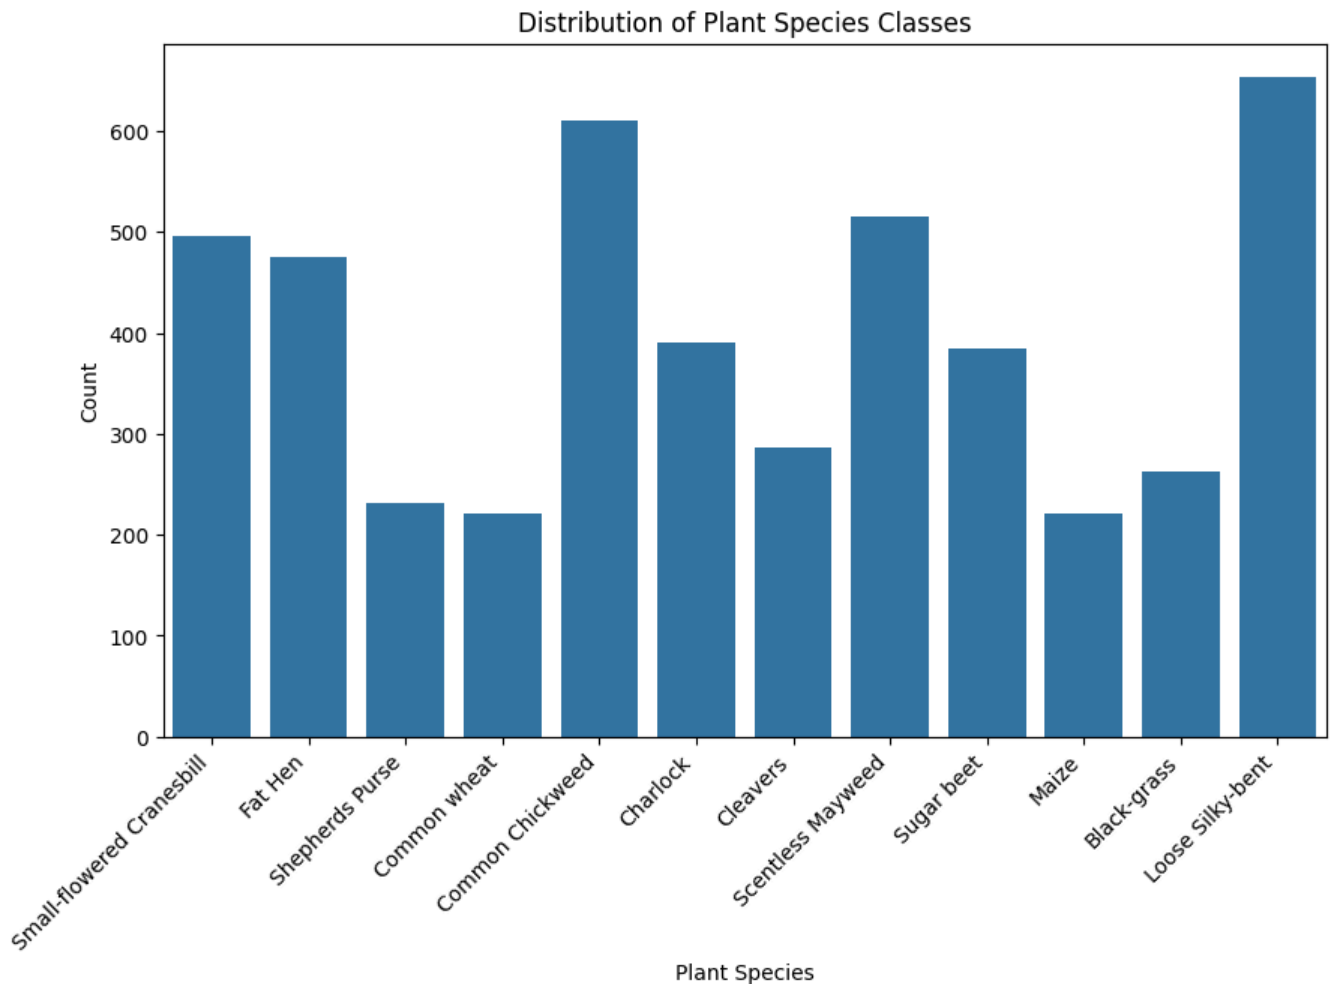
Images shape: (4750, 128, 128, 3)
Labels head:

	Label
0	Small-flowered Cranesbill
1	Small-flowered Cranesbill
2	Small-flowered Cranesbill
3	Small-flowered Cranesbill
4	Small-flowered Cranesbill

```
1 # Display column names to check for issues
2 print("Column names in labels DataFrame:", labels.columns)
3
```

➡ Column names in labels DataFrame: Index(['Label'], dtype='object')

```
1 # Plot the distribution of plant species classes with the correct column name
2 plt.figure(figsize=(10, 6))
3 sns.countplot(data=labels, x='Label') # Using 'Label' instead of 'species'
4 plt.xticks(rotation=45, ha='right') # Rotate x labels for readability
5 plt.title('Distribution of Plant Species Classes')
6 plt.xlabel('Plant Species')
7 plt.ylabel('Count')
8 plt.show()
9
```



2. Data Preprocessing

```
1 # Ensure we have at least 12 unique species
2 unique_species = labels['Label'].unique()
3 if len(unique_species) < 12:
4     print("Not enough species to display a 3x4 grid.")
5 else:
6     # Select one sample image per species (12 in total)
7     samples_per_class = {}
8     for species in unique_species[:12]: # Limit to the first 12 unique species
9         sample_index = labels[labels['Label'] == species].index[0]
10        sample_image = images[sample_index]
11
12        # Reshape or normalize if needed (check your image's original dimensions)
13        if sample_image.shape[-1] != 3: # Assuming RGB images
14            sample_image = np.repeat(sample_image[..., np.newaxis], 3, axis=-1)
15
16        samples_per_class[species] = sample_image
17
18    # Plotting the images in a 3x4 grid
19    fig, axes = plt.subplots(3, 4, figsize=(12, 9))
20    fig.suptitle('Sample Images from Each Plant Species', fontsize=16)
21
22    for ax, (species, image) in zip(axes.flat, samples_per_class.items()):
23        ax.imshow(image)
24        ax.set_title(species)
25        ax.axis('off')
26
27    plt.tight_layout()
28    plt.subplots_adjust(top=0.88)
29    plt.show()
30
```



Sample Images from Each Plant Species

Small-flowered Cranesbill




Fat Hen



Shepherds Purse



Common wheat



Common Chickweed



Charlock




Cleavers



Scentless Mayweed




Sugar beet




Maize



Black-grass



Loose Silky-bent



```
1 import cv2
2
```

```

3 # Resize images to 64x64
4 resized_images = np.array([cv2.resize(image, (64, 64)) for image in images])
5 print("Resized images shape:", resized_images.shape)
6

```

➡ Resized images shape: (4750, 64, 64, 3)

```

1 from sklearn.preprocessing import LabelEncoder
2 from tensorflow.keras.utils import to_categorical
3
4 # Encode labels as integers
5 label_encoder = LabelEncoder()
6 encoded_labels = label_encoder.fit_transform(labels['Label'])
7
8 # Convert integer-encoded labels to one-hot encoded labels
9 one_hot_labels = to_categorical(encoded_labels)
10 print("One-hot encoded labels shape:", one_hot_labels.shape)
11

```

➡ One-hot encoded labels shape: (4750, 12)

```

1 # Normalize pixel values to range [0, 1]
2 normalized_images = resized_images / 255.0
3 print("Normalized images shape:", normalized_images.shape)
4

```

➡ Normalized images shape: (4750, 64, 64, 3)

3. Model 1: Basic CNN Model

```


1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
3 from tensorflow.keras.regularizers import l2
4 from tensorflow.keras.callbacks import EarlyStopping
5
6 # Define the Basic CNN architecture with a deeper dense network
7 model_1 = Sequential([
8     Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
9     MaxPooling2D((2, 2)),
10
11     Conv2D(64, (3, 3), activation='relu'),
12     MaxPooling2D((2, 2)),
13
14     Conv2D(128, (3, 3), activation='relu'),
15     MaxPooling2D((2, 2)),
16
17     Flatten(),
18

```

```

19 # Dense Layers with regularization
20 Dense(256, activation='relu', kernel_regularizer=l2(0.001)),
21 Dropout(0.4),
22 Dense(128, activation='relu', kernel_regularizer=l2(0.001)),
23 Dropout(0.4),
24
25 # Output layer for 12 classes
26 Dense(12, activation='softmax')
27 ])
28
29 # Compile the model
30 model_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
31
32 # Display the model summary
33 model_1.summary()
34

```

 /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107:
 super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 256)	1,179,904
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32,896
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 12)	1,548



 Total params: 1,307,596 (4.00 MB)

```

1 # Define early stopping
2 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
3

```

```

4 # Split the data into training and validation sets
5 from sklearn.model_selection import train_test_split
6 X_train, X_val, y_train, y_val = train_test_split(normalized_images, one_hot_labels, test
7
8 # Train the model
9 history = model_1.fit(X_train, y_train, epochs=5, validation_data=(X_val, y_val),
10                       callbacks=[early_stopping], batch_size=32)
11

```



Epoch 1/5

119/119 ————— **20s** 100ms/step - accuracy: 0.1321 - loss: 2.7617 - val_accu

Epoch 2/5

119/119 ————— **1s** 7ms/step - accuracy: 0.3361 - loss: 1.9809 - val_accurac

Epoch 3/5

119/119 ————— **1s** 6ms/step - accuracy: 0.4158 - loss: 1.7201 - val_accurac

Epoch 4/5

119/119 ————— **1s** 6ms/step - accuracy: 0.4467 - loss: 1.6166 - val_accurac

Epoch 5/5



119/119 ————— **1s** 7ms/step - accuracy: 0.5234 - loss: 1.4301 - val_accurac



```

1 import matplotlib.pyplot as plt
2
3 # Plot accuracy and loss curves
4 plt.figure(figsize=(12, 4))
5
6 # Accuracy plot
7 plt.subplot(1, 2, 1)
8 plt.plot(history.history['accuracy'], label='Training Accuracy')
9 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
10 plt.xlabel('Epoch')
11 plt.ylabel('Accuracy')
12 plt.legend()
13 plt.title('Model Accuracy')
14
15 # Loss plot
16 plt.subplot(1, 2, 2)
17 plt.plot(history.history['loss'], label='Training Loss')
18 plt.plot(history.history['val_loss'], label='Validation Loss')
19 plt.xlabel('Epoch')
20 plt.ylabel('Loss')
21 plt.legend()
22 plt.title('Model Loss')
23
24 plt.show()
25

```




```
1 from sklearn.metrics import confusion_matrix, classification_report
2 import seaborn as sns
3
4 # Make predictions on the validation set
5 y_pred = model_1.predict(X_val)
6 y_pred_classes = np.argmax(y_pred, axis=1)
7 y_true_classes = np.argmax(y_val, axis=1)
8
9 # Generate confusion matrix
10 conf_matrix = confusion_matrix(y_true_classes, y_pred_classes)
11
12 # Plot the confusion matrix
13 plt.figure(figsize=(10, 8))
14 sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.clas
15             yticklabels=label_encoder.classes_)
16 plt.xlabel('Predicted Labels')
17 plt.ylabel('True Labels')
18 plt.title('Confusion Matrix')
19 plt.show()
20
21 # Print classification report
22 print("Classification Report:\n", classification_report(y_true_classes, y_pred_classes, t
23
```




30/30

1s 10ms/step



Classification Report:

	precision	recall	f1-score	support
Black-grass	0.00	0.00	0.00	47
Charlock	0.84	0.68	0.75	76
Cleavers	0.77	0.39	0.52	62
Common Chickweed	0.59	0.93	0.72	102
Common wheat	0.00	0.00	0.00	38
Fat Hen	0.45	0.50	0.47	94
Loose Silky-bent	0.59	0.94	0.72	141
Maize	0.96	0.52	0.68	42
Scentless Mayweed	0.53	0.75	0.62	105
Shepherds Purse	0.36	0.07	0.12	56
Small-flowered Cranesbill	0.80	0.82	0.81	114
Sugar beet	0.50	0.44	0.47	73
accuracy			0.61	950
macro avg	0.53	0.50	0.49	950
weighted avg	0.57	0.61	0.56	950

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531: Unde
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531: Unde
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```


```
1 # Plot a few test samples with predicted and true labels
2 fig, axes = plt.subplots(3, 3, figsize=(10, 10))
3 axes = axes.flatten()
4
5 for i in range(9):
6     img = X_val[i]
7     true_label = label_encoder.classes_[y_true_classes[i]]
8     pred_label = label_encoder.classes_[y_pred_classes[i]]
9
10    axes[i].imshow(img)
11    axes[i].set_title(f"True: {true_label}\nPred: {pred_label}")
12    axes[i].axis('off')
13
14 plt.tight_layout()
15 plt.show()
16
```



True: Maize
Pred: Common Chickweed



True: Shepherds Purse
Pred: Small-flowered Cranesbill



True: Small-flowered Cranesbill
Pred: Small-flowered Cranesbill



True: Charlock
Pred: Scentless Mayweed



True: Common Chickweed
Pred: Common Chickweed




True: Loose Silky-bent
Pred: Loose Silky-bent



True: Maize
Pred: Maize



True: Sugar beet
Pred: Sugar beet



True: Common wheat
Pred: Loose Silky-bent



4. Model 2: Enhanced CNN with Data Augmentation and Regularization

```
1 #Clearing backend
2 from tensorflow.keras import backend
3 backend.clear_session()
```

```
1 from tensorflow.keras.preprocessing.image import ImageDataGenerator
2
3 # Define data augmentation
4 datagen = ImageDataGenerator(
5     rotation_range=20,      # Rotate images up to 20 degrees
6     width_shift_range=0.1,  # Shift width by up to 10%
7     height_shift_range=0.1, # Shift height by up to 10%
8     zoom_range=0.2,        # Zoom in by up to 20%
9     horizontal_flip=True    # Randomly flip images horizontally
10 )
11
12 # Fit the data generator on the training data
13 datagen.fit(X_train)
14
```

```
1 from tensorflow.keras.preprocessing.image import ImageDataGenerator
2 from tensorflow.keras.layers import SpatialDropout2D
3
4 # Define the Enhanced CNN architecture with deeper dense layers
5 model_2 = Sequential([
6     Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
7     BatchNormalization(),
8     MaxPooling2D((2, 2)),
9     SpatialDropout2D(0.2),
10
11     Conv2D(64, (3, 3), activation='relu'),
12     BatchNormalization(),
13     MaxPooling2D((2, 2)),
14     SpatialDropout2D(0.3),
15
16     Conv2D(128, (3, 3), activation='relu'),
17     BatchNormalization(),
18     MaxPooling2D((2, 2)),
19     SpatialDropout2D(0.4),
20
21     Flatten(),
22
23     # Deeper Dense Layers with regularization
24     Dense(256, activation='relu', kernel_regularizer=l2(0.001)),
25     Dropout(0.4),
```

```
26     Dense(128, activation='relu', kernel_regularizer=l2(0.001)),
27     Dropout(0.4),
28
29     # Output layer for 12 classes
30     Dense(12, activation='softmax')
31 ])
32
33 # Compile the model
34 model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
35
36 # Display the model summary
37 model_2.summary()
38
```

```
➡ /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107:
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None , 62, 62, 32)	896
batch_normalization (BatchNormalization)	(None , 62, 62, 32)	128
max_pooling2d (MaxPooling2D)	(None , 31, 31, 32)	0
spatial_dropout2d (SpatialDropout2D)	(None , 31, 31, 32)	0
conv2d_1 (Conv2D)	(None , 29, 29, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None , 29, 29, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None , 14, 14, 64)	0
spatial_dropout2d_1 (SpatialDropout2D)	(None , 14, 14, 64)	0
conv2d_2 (Conv2D)	(None , 12, 12, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None , 12, 12, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None , 6, 6, 128)	0
spatial_dropout2d_2 (SpatialDropout2D)	(None , 6, 6, 128)	0
flatten (Flatten)	(None , 4608)	0
dense (Dense)	(None , 256)	1,179,904
dropout (Dropout)	(None , 256)	0
dense_1 (Dense)	(None , 128)	32,896
dropout_1 (Dropout)	(None , 128)	0
dense_2 (Dense)	(None , 12)	1,548

[illegible]

```
7 epochs=5,  
8 callbacks=[early_stopping])  
9
```



Epoch 1/5

/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:119: self._warn_if_super_not_called()

119/119 ————— **24s** 130ms/step - accuracy: 0.1442 - loss: 4.1249 - val_accu

Epoch 2/5

119/119 ————— **6s** 45ms/step - accuracy: 0.2509 - loss: 2.9522 - val_accu

Epoch 3/5

119/119 ————— **11s** 48ms/step - accuracy: 0.3222 - loss: 2.7471 - val_accu

Epoch 4/5



119/119 ————— **7s** 60ms/step - accuracy: 0.4009 - loss: 2.4982 - val_accu

Epoch 5/5

119/119 ————— **8s** 42ms/step - accuracy: 0.4002 - loss: 2.4475 - val_accu



```
1 # Plot accuracy and loss curves for the enhanced model  
2 plt.figure(figsize=(12, 4))  
3  
4 # Accuracy plot  
5 plt.subplot(1, 2, 1)  
6 plt.plot(history_enhanced.history['accuracy'], label='Training Accuracy')  
7 plt.plot(history_enhanced.history['val_accuracy'], label='Validation Accuracy')  
8 plt.xlabel('Epoch')  
9 plt.ylabel('Accuracy')  
10 plt.legend()  
11 plt.title('Enhanced Model Accuracy')  
12  
13 # Loss plot  
14 plt.subplot(1, 2, 2)  
15 plt.plot(history_enhanced.history['loss'], label='Training Loss')  
16 plt.plot(history_enhanced.history['val_loss'], label='Validation Loss')  
17 plt.xlabel('Epoch')  
18 plt.ylabel('Loss')  
19 plt.legend()  
20 plt.title('Enhanced Model Loss')  
21  
22 plt.show()  
23
```



```
1 # Make predictions on the validation set
2 y_pred_enhanced = model_2.predict(X_val)
3 y_pred_classes_enhanced = np.argmax(y_pred_enhanced, axis=1)
4 y_true_classes_enhanced = np.argmax(y_val, axis=1)
5
6 # Generate confusion matrix
7 conf_matrix_enhanced = confusion_matrix(y_true_classes_enhanced, y_pred_classes_enhanced)
8
9 # Plot the confusion matrix
10 plt.figure(figsize=(10, 8))
11 sns.heatmap(conf_matrix_enhanced, annot=True, fmt='d', cmap='Blues',
12             xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
13 plt.xlabel('Predicted Labels')
14 plt.ylabel('True Labels')
15 plt.title('Enhanced Model Confusion Matrix')
16 plt.show()
17
18 # Print classification report
19 print("Enhanced Model Classification Report:\n",
20       classification_report(y_true_classes_enhanced, y_pred_classes_enhanced, target_name
21
```




Enhanced Model Classification Report:				
	precision	recall	f1-score	support
Black-grass	0.26	0.38	0.31	47
Charlock	0.35	0.88	0.50	76
Cleavers	0.39	0.73	0.51	62
Common Chickweed	0.68	0.35	0.46	102
Common wheat	0.18	0.34	0.23	38
Fat Hen	0.32	0.07	0.12	94
Loose Silky-bent	0.57	0.74	0.64	141
Maize	0.00	0.00	0.00	42
Scentless Mayweed	0.25	0.31	0.28	105
Shepherds Purse	0.75	0.05	0.10	56
Small-flowered Cranesbill	0.83	0.56	0.67	114
Sugar beet	0.92	0.33	0.48	73
accuracy			0.44	950
macro avg	0.46	0.40	0.36	950
weighted avg	0.50	0.44	0.41	950

```
1 # Plot a few test samples with predicted and true labels
2 fig, axes = plt.subplots(3, 3, figsize=(10, 10))
3 axes = axes.flatten()
4
5 for i in range(9):
6     img = X_val[i]
7     true_label = label_encoder.classes_[y_true_classes_enhanced[i]]
8     pred_label = label_encoder.classes_[y_pred_classes_enhanced[i]]
9
10    axes[i].imshow(img)
11    axes[i].set_title(f"True: {true_label}\nPred: {pred_label}")
12    axes[i].axis('off')
13
14 plt.tight_layout()
15 plt.show()
16
```



True: Maize
Pred: Scentless Mayweed



True: Shepherds Purse
Pred: Charlock



True: Small-flowered Cranesbill
Pred: Small-flowered Cranesbill



True: Charlock
Pred: Charlock



True: Common Chickweed
Pred: Common Chickweed




True: Loose Silky-bent
Pred: Black-grass



True: Maize
Pred: Common Chickweed



True: Sugar beet
Pred: Sugar beet



True: Common wheat
Pred: Loose Silky-bent



5. Model 3: Transfer Learning with Pre-trained Models

```
1 !pip install keras
2 !pip install keras-tuner
3 from tensorflow.keras.applications import VGG16, ResNet50
4 from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout, BatchNormaliz
5 from tensorflow.keras.models import Model
6 from tensorflow.keras.optimizers import Adam
7 from keras_tuner import HyperModel, RandomSearch # This import should now work
8 from tensorflow.keras.regularizers import l2
```

➡ Requirement already satisfied: keras in /usr/local/lib/python3.10/dist-packages (3.4.1)
Requirement already satisfied: absl-py in /usr/local/lib/python3.10/dist-packages (from keras) (0.15.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from keras) (1.24.3)
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras) (13.7.1)
Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras) (0.0.8)
Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages (from keras) (3.10.0)
Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras) (0.9.2)
Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.10/dist-packages (from keras) (0.2.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from keras) (23.1)
Requirement already satisfied: typing-extensions>=4.5.0 in /usr/local/lib/python3.10/dist-packages (from keras) (4.5.0)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich) (2.17.2)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py) (0.1.2)
Collecting keras-tuner
 Downloading keras_tuner-1.4.7-py3-none-any.whl.metadata (5.4 kB)
Requirement already satisfied: keras in /usr/local/lib/python3.10/dist-packages (from keras-tuner) (3.4.1)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from keras-tuner) (23.1)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from keras-tuner) (2.31.0)
Collecting kt-legacy (from keras-tuner)
 Downloading kt_legacy-1.0.5-py3-none-any.whl.metadata (221 bytes)
Requirement already satisfied: absl-py in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (0.15.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (1.24.3)
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (13.7.1)
Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (0.0.8)
Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (3.10.0)
Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (0.9.2)
Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.10/dist-packages (from kt-legacy) (0.2.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests) (2024.2.2)
Requirement already satisfied: typing-extensions>=4.5.0 in /usr/local/lib/python3.10/dist-packages (from requests) (4.5.0)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich) (2.17.2)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py) (0.1.2)
Downloading keras_tuner-1.4.7-py3-none-any.whl (129 kB)
129.1/129.1 kB 5.2 MB/s eta 0:00:00
Downloading kt_legacy-1.0.5-py3-none-any.whl (9.6 kB)
Installing collected packages: kt-legacy, keras-tuner
Successfully installed keras-tuner-1.4.7 kt-legacy-1.0.5

```

1 class TransferLearningHyperModel(HyperModel):
2     def __init__(self, base_model):
3         self.base_model = base_model
4
5     def build(self, hp):
6         self.base_model.trainable = False
7
8         # Add layers to the base model
9         x = self.base_model.output
10        x = GlobalAveragePooling2D()(x)
11
12        # Add dense layers with regularization
13        for i in range(hp.Int('num_dense_layers', 2, 3)):
14            units = hp.Int(f'units_{i}', min_value=64, max_value=256, step=32)
15            x = Dense(units, activation='relu', kernel_regularizer=l2(0.001))(x)
16            x = BatchNormalization()(x)
17            x = Dropout(hp.Float('dropout_rate', 0.3, 0.5))(x)
18
19        # Final output layer
20        output = Dense(12, activation='softmax')(x)
21        model = Model(inputs=self.base_model.input, outputs=output)
22
23        # Compile model
24        learning_rate = hp.Choice('learning_rate', [1e-4, 1e-3, 5e-4])
25        model.compile(optimizer=Adam(learning_rate=learning_rate),
26                      loss='categorical_crossentropy',
27                      metrics=['accuracy'])
28        return model
29

```

```


1 # VGG16 Transfer Learning with Tuner
2 vgg16_base = VGG16(weights='imagenet', include_top=False, input_shape=(64, 64, 3))
3 vgg16_hypermodel = TransferLearningHyperModel(vgg16_base)
4
5 vgg16_tuner = RandomSearch(
6     vgg16_hypermodel,
7     objective='val_accuracy',
8     max_trials=5,
9     executions_per_trial=2,
10    directory='vgg16_tuner',
11    project_name='vgg16_tuning'
12 )
13
14 vgg16_tuner.search(X_train, y_train, epochs=5, validation_data=(X_val, y_val))
15 best_vgg16_model = vgg16_tuner.get_best_models(num_models=1)[0]
16
17 # ResNet50 Transfer Learning with Tuner
18 resnet50_base = ResNet50(weights='imagenet', include_top=False, input_shape=(64, 64, 3))

```

```

19 resnet50_hypermodel = TransferLearningHyperModel(resnet50_base)
20
21 resnet50_tuner = RandomSearch(
22     resnet50_hypermodel,
23     objective='val_accuracy',
24     max_trials=5,
25     executions_per_trial=2,
26     directory='resnet50_tuner',
27     project_name='resnet50_tuning'
28 )
29
30 resnet50_tuner.search(X_train, y_train, epochs=5, validation_data=(X_val, y_val))
31 best_resnet50_model = resnet50_tuner.get_best_models(num_models=1)[0]
32

```

 Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/vgg16/58889256/58889256> **0s** 0us/step
 Reloading Tuner from vgg16_tuner/vgg16_tuning/tuner0.json
 /usr/local/lib/python3.10/dist-packages/keras/src/saving/saving_lib.py:576: UserWarning:
 saveable.load_own_variables(weights_store.get(inner_path))
 Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/resnet50/94765736/94765736> **1s** 0us/step
 Reloading Tuner from resnet50_tuner/resnet50_tuning/tuner0.json

```

1 # For the VGG16 transfer learning model
2 best_vgg16_model.summary()

```



Model: "functional"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 64, 64, 3)	0
block1_conv1 (Conv2D)	(None, 64, 64, 64)	1,792
block1_conv2 (Conv2D)	(None, 64, 64, 64)	36,928
block1_pool1 (MaxPooling2D)	(None, 32, 32, 64)	0
block2_conv1 (Conv2D)	(None, 32, 32, 128)	73,856
block2_conv2 (Conv2D)	(None, 32, 32, 128)	147,584
block2_pool1 (MaxPooling2D)	(None, 16, 16, 128)	0
block3_conv1 (Conv2D)	(None, 16, 16, 256)	295,168
block3_conv2 (Conv2D)	(None, 16, 16, 256)	590,080
block3_conv3 (Conv2D)	(None, 16, 16, 256)	590,080
block3_pool1 (MaxPooling2D)	(None, 8, 8, 256)	0
block4_conv1 (Conv2D)	(None, 8, 8, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 8, 8, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 8, 8, 512)	2,359,808
block4_pool1 (MaxPooling2D)	(None, 4, 4, 512)	0
block5_conv1 (Conv2D)	(None, 4, 4, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 4, 4, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 4, 4, 512)	2,359,808
block5_pool1 (MaxPooling2D)	(None, 2, 2, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense (Dense)	(None, 256)	131,328
batch_normalization (BatchNormalization)	(None, 256)	1,024
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 256)	65,792
batch_normalization_1 (BatchNormalization)	(None, 256)	1,024

dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 12)	3,084

```
1 # For the ResNet50 transfer learning model
2 best_resnet50_model.summary()
3
```




Model: "functional"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 64, 64, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 70, 70, 3)	0	input_layer[0][0]
conv1_conv (Conv2D)	(None, 32, 32, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 32, 32, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 32, 32, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 34, 34, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 16, 16, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 16, 16, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv2_block1_1_conv
conv2_block1_1_relu (Activation)	(None, 16, 16, 64)	0	conv2_block1_1_bn
conv2_block1_2_conv (Conv2D)	(None, 16, 16, 64)	36,928	conv2_block1_1_relu
conv2_block1_2_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv2_block1_2_conv
conv2_block1_2_relu (Activation)	(None, 16, 16, 64)	0	conv2_block1_2_bn
conv2_block1_0_conv (Conv2D)	(None, 16, 16, 256)	16,640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 16, 16, 256)	16,640	conv2_block1_2_relu
conv2_block1_0_bn (BatchNormalization)	(None, 16, 16, 256)	1,024	conv2_block1_0_conv
conv2_block1_3_bn (BatchNormalization)	(None, 16, 16, 256)	1,024	conv2_block1_3_conv
conv2_block1_add (Add)	(None, 16, 16, 256)	0	conv2_block1_0_bn conv2_block1_3_bn
conv2_block1_out (Activation)	(None, 16, 16, 256)	0	conv2_block1_add
conv2_block2_1_conv (Conv2D)	(None, 16, 16, 64)	16,448	conv2_block1_out

conv2_block2_1_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv2_block2_1_c
conv2_block2_1_relu (Activation)	(None, 16, 16, 64)	0	conv2_block2_1_t
conv2_block2_2_conv (Conv2D)	(None, 16, 16, 64)	36,928	conv2_block2_1_r
conv2_block2_2_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv2_block2_2_c
conv2_block2_2_relu (Activation)	(None, 16, 16, 64)	0	conv2_block2_2_t
conv2_block2_3_conv (Conv2D)	(None, 16, 16, 256)	16,640	conv2_block2_2_r
conv2_block2_3_bn (BatchNormalization)	(None, 16, 16, 256)	1,024	conv2_block2_3_c
conv2_block2_add (Add)	(None, 16, 16, 256)	0	conv2_block1_out conv2_block2_3_t
conv2_block2_out (Activation)	(None, 16, 16, 256)	0	conv2_block2_add
conv2_block3_1_conv (Conv2D)	(None, 16, 16, 64)	16,448	conv2_block2_out
conv2_block3_1_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv2_block3_1_c
conv2_block3_1_relu (Activation)	(None, 16, 16, 64)	0	conv2_block3_1_t
conv2_block3_2_conv (Conv2D)	(None, 16, 16, 64)	36,928	conv2_block3_1_r
conv2_block3_2_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv2_block3_2_c
conv2_block3_2_relu (Activation)	(None, 16, 16, 64)	0	conv2_block3_2_t
conv2_block3_3_conv (Conv2D)	(None, 16, 16, 256)	16,640	conv2_block3_2_r
conv2_block3_3_bn (BatchNormalization)	(None, 16, 16, 256)	1,024	conv2_block3_3_c
conv2_block3_add (Add)	(None, 16, 16, 256)	0	conv2_block2_out conv2_block3_3_t
conv2_block3_out (Activation)	(None, 16, 16, 256)	0	conv2_block3_add

conv3_block1_1_conv (Conv2D)	(None, 8, 8, 128)	32,896	conv2_block3_out
conv3_block1_1_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block1_1_c
conv3_block1_1_relu (Activation)	(None, 8, 8, 128)	0	conv3_block1_1_t
conv3_block1_2_conv (Conv2D)	(None, 8, 8, 128)	147,584	conv3_block1_1_r
conv3_block1_2_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block1_2_c
conv3_block1_2_relu (Activation)	(None, 8, 8, 128)	0	conv3_block1_2_t
conv3_block1_0_conv (Conv2D)	(None, 8, 8, 512)	131,584	conv2_block3_out
conv3_block1_3_conv (Conv2D)	(None, 8, 8, 512)	66,048	conv3_block1_2_r
conv3_block1_0_bn (BatchNormalization)	(None, 8, 8, 512)	2,048	conv3_block1_0_c
conv3_block1_3_bn (BatchNormalization)	(None, 8, 8, 512)	2,048	conv3_block1_3_c
conv3_block1_add (Add)	(None, 8, 8, 512)	0	conv3_block1_0_t conv3_block1_3_t
conv3_block1_out (Activation)	(None, 8, 8, 512)	0	conv3_block1_add
conv3_block2_1_conv (Conv2D)	(None, 8, 8, 128)	65,664	conv3_block1_out
conv3_block2_1_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block2_1_c
conv3_block2_1_relu (Activation)	(None, 8, 8, 128)	0	conv3_block2_1_t
conv3_block2_2_conv (Conv2D)	(None, 8, 8, 128)	147,584	conv3_block2_1_r
conv3_block2_2_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block2_2_c
conv3_block2_2_relu (Activation)	(None, 8, 8, 128)	0	conv3_block2_2_t
conv3_block2_3_conv (Conv2D)	(None, 8, 8, 512)	66,048	conv3_block2_2_r
conv3_block2_3_bn	(None, 8, 8, 512)	2,048	conv3_block2_3_c

(BatchNormalization)	(None, 8, 8, 512)	0	
conv3_block2_add (Add)	(None, 8, 8, 512)	0	conv3_block1_out conv3_block2_3_t
conv3_block2_out (Activation)	(None, 8, 8, 512)	0	conv3_block2_add
conv3_block3_1_conv (Conv2D)	(None, 8, 8, 128)	65,664	conv3_block2_out
conv3_block3_1_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block3_1_c
conv3_block3_1_relu (Activation)	(None, 8, 8, 128)	0	conv3_block3_1_t
conv3_block3_2_conv (Conv2D)	(None, 8, 8, 128)	147,584	conv3_block3_1_r
conv3_block3_2_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block3_2_c
conv3_block3_2_relu (Activation)	(None, 8, 8, 128)	0	conv3_block3_2_t
conv3_block3_3_conv (Conv2D)	(None, 8, 8, 512)	66,048	conv3_block3_2_r
conv3_block3_3_bn (BatchNormalization)	(None, 8, 8, 512)	2,048	conv3_block3_3_c
conv3_block3_add (Add)	(None, 8, 8, 512)	0	conv3_block2_out conv3_block3_3_t
conv3_block3_out (Activation)	(None, 8, 8, 512)	0	conv3_block3_add
conv3_block4_1_conv (Conv2D)	(None, 8, 8, 128)	65,664	conv3_block3_out
conv3_block4_1_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block4_1_c
conv3_block4_1_relu (Activation)	(None, 8, 8, 128)	0	conv3_block4_1_t
conv3_block4_2_conv (Conv2D)	(None, 8, 8, 128)	147,584	conv3_block4_1_r
conv3_block4_2_bn (BatchNormalization)	(None, 8, 8, 128)	512	conv3_block4_2_c
conv3_block4_2_relu (Activation)	(None, 8, 8, 128)	0	conv3_block4_2_t
conv3_block4_3_conv (Conv2D)	(None, 8, 8, 512)	66,048	conv3_block4_2_r

(Conv2D)			
conv3_block4_3_bn (BatchNormalization)	(None, 8, 8, 512)	2,048	conv3_block4_3_c
conv3_block4_add (Add)	(None, 8, 8, 512)	0	conv3_block3_out conv3_block4_3_t
conv3_block4_out (Activation)	(None, 8, 8, 512)	0	conv3_block4_add
conv4_block1_1_conv (Conv2D)	(None, 4, 4, 256)	131,328	conv3_block4_out
conv4_block1_1_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block1_1_c
conv4_block1_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block1_1_t
conv4_block1_2_conv (Conv2D)	(None, 4, 4, 256)	590,080	conv4_block1_1_r
conv4_block1_2_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block1_2_c
conv4_block1_2_relu (Activation)	(None, 4, 4, 256)	0	conv4_block1_2_t
conv4_block1_0_conv (Conv2D)	(None, 4, 4, 1024)	525,312	conv3_block4_out
conv4_block1_3_conv (Conv2D)	(None, 4, 4, 1024)	263,168	conv4_block1_2_r
conv4_block1_0_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block1_0_c
conv4_block1_3_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block1_3_c
conv4_block1_add (Add)	(None, 4, 4, 1024)	0	conv4_block1_0_t conv4_block1_3_t
conv4_block1_out (Activation)	(None, 4, 4, 1024)	0	conv4_block1_add
conv4_block2_1_conv (Conv2D)	(None, 4, 4, 256)	262,400	conv4_block1_out
conv4_block2_1_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block2_1_c
conv4_block2_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block2_1_t
conv4_block2_2_conv (Conv2D)	(None, 4, 4, 256)	590,080	conv4_block2_1_r

conv4_block2_2_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block2_2_c
conv4_block2_2_relu (Activation)	(None, 4, 4, 256)	0	conv4_block2_2_t
conv4_block2_3_conv (Conv2D)	(None, 4, 4, 1024)	263,168	conv4_block2_2_r
conv4_block2_3_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block2_3_c
conv4_block2_add (Add)	(None, 4, 4, 1024)	0	conv4_block1_out conv4_block2_3_t
conv4_block2_out (Activation)	(None, 4, 4, 1024)	0	conv4_block2_add
conv4_block3_1_conv (Conv2D)	(None, 4, 4, 256)	262,400	conv4_block2_out
conv4_block3_1_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block3_1_c
conv4_block3_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block3_1_t
conv4_block3_2_conv (Conv2D)	(None, 4, 4, 256)	590,080	conv4_block3_1_r
conv4_block3_2_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block3_2_c
conv4_block3_2_relu (Activation)	(None, 4, 4, 256)	0	conv4_block3_2_t
conv4_block3_3_conv (Conv2D)	(None, 4, 4, 1024)	263,168	conv4_block3_2_r
conv4_block3_3_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block3_3_c
conv4_block3_add (Add)	(None, 4, 4, 1024)	0	conv4_block2_out conv4_block3_3_t
conv4_block3_out (Activation)	(None, 4, 4, 1024)	0	conv4_block3_add
conv4_block4_1_conv (Conv2D)	(None, 4, 4, 256)	262,400	conv4_block3_out
conv4_block4_1_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block4_1_c
conv4_block4_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block4_1_t

conv4_block4_2_conv (Conv2D)	(None, 4, 4, 256)	590,080	conv4_block4_1_r
conv4_block4_2_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block4_2_c
conv4_block4_2_relu (Activation)	(None, 4, 4, 256)	0	conv4_block4_2_t
conv4_block4_3_conv (Conv2D)	(None, 4, 4, 1024)	263,168	conv4_block4_2_r
conv4_block4_3_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block4_3_c
conv4_block4_add (Add)	(None, 4, 4, 1024)	0	conv4_block3_out conv4_block4_3_t
conv4_block4_out (Activation)	(None, 4, 4, 1024)	0	conv4_block4_add
conv4_block5_1_conv (Conv2D)	(None, 4, 4, 256)	262,400	conv4_block4_out
conv4_block5_1_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block5_1_c
conv4_block5_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block5_1_t
conv4_block5_2_conv (Conv2D)	(None, 4, 4, 256)	590,080	conv4_block5_1_r
conv4_block5_2_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block5_2_c
conv4_block5_2_relu (Activation)	(None, 4, 4, 256)	0	conv4_block5_2_t
conv4_block5_3_conv (Conv2D)	(None, 4, 4, 1024)	263,168	conv4_block5_2_r
conv4_block5_3_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block5_3_c
conv4_block5_add (Add)	(None, 4, 4, 1024)	0	conv4_block4_out conv4_block5_3_t
conv4_block5_out (Activation)	(None, 4, 4, 1024)	0	conv4_block5_add
conv4_block6_1_conv (Conv2D)	(None, 4, 4, 256)	262,400	conv4_block5_out
conv4_block6_1_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block6_1_c
conv4_block6_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block6_1_t

conv4_block6_1_relu (Activation)	(None, 4, 4, 256)	0	conv4_block6_1_t
conv4_block6_2_conv (Conv2D)	(None, 4, 4, 256)	590,080	conv4_block6_1_r
conv4_block6_2_bn (BatchNormalization)	(None, 4, 4, 256)	1,024	conv4_block6_2_c
conv4_block6_2_relu (Activation)	(None, 4, 4, 256)	0	conv4_block6_2_t
conv4_block6_3_conv (Conv2D)	(None, 4, 4, 1024)	263,168	conv4_block6_2_r
conv4_block6_3_bn (BatchNormalization)	(None, 4, 4, 1024)	4,096	conv4_block6_3_c
conv4_block6_add (Add)	(None, 4, 4, 1024)	0	conv4_block5_out conv4_block6_3_t
conv4_block6_out (Activation)	(None, 4, 4, 1024)	0	conv4_block6_add
conv5_block1_1_conv (Conv2D)	(None, 2, 2, 512)	524,800	conv4_block6_out
conv5_block1_1_bn (BatchNormalization)	(None, 2, 2, 512)	2,048	conv5_block1_1_c
conv5_block1_1_relu (Activation)	(None, 2, 2, 512)	0	conv5_block1_1_t
conv5_block1_2_conv (Conv2D)	(None, 2, 2, 512)	2,359,808	conv5_block1_1_r
conv5_block1_2_bn (BatchNormalization)	(None, 2, 2, 512)	2,048	conv5_block1_2_c
conv5_block1_2_relu (Activation)	(None, 2, 2, 512)	0	conv5_block1_2_t
conv5_block1_0_conv (Conv2D)	(None, 2, 2, 2048)	2,099,200	conv4_block6_out
conv5_block1_3_conv (Conv2D)	(None, 2, 2, 2048)	1,050,624	conv5_block1_2_r
conv5_block1_0_bn (BatchNormalization)	(None, 2, 2, 2048)	8,192	conv5_block1_0_c
conv5_block1_3_bn (BatchNormalization)	(None, 2, 2, 2048)	8,192	conv5_block1_3_c
conv5_block1_add (Add)	(None, 2, 2, 2048)	0	conv5_block1_0_t conv5_block1_3_t
conv5_block1_out	(None, 2, 2, 2048)	0	conv5_block1_add

(Activation)			
conv5_block2_1_conv (Conv2D)	(None, 2, 2, 512)	1,049,088	conv5_block1_out
conv5_block2_1_bn (BatchNormalization)	(None, 2, 2, 512)	2,048	conv5_block2_1_c
conv5_block2_1_relu (Activation)	(None, 2, 2, 512)	0	conv5_block2_1_t
conv5_block2_2_conv (Conv2D)	(None, 2, 2, 512)	2,359,808	conv5_block2_1_r
conv5_block2_2_bn (BatchNormalization)	(None, 2, 2, 512)	2,048	conv5_block2_2_c
conv5_block2_2_relu (Activation)	(None, 2, 2, 512)	0	conv5_block2_2_t
conv5_block2_3_conv (Conv2D)	(None, 2, 2, 2048)	1,050,624	conv5_block2_2_r
conv5_block2_3_bn (BatchNormalization)	(None, 2, 2, 2048)	8,192	conv5_block2_3_c
conv5_block2_add (Add)	(None, 2, 2, 2048)	0	conv5_block1_out conv5_block2_3_t
conv5_block2_out (Activation)	(None, 2, 2, 2048)	0	conv5_block2_add
conv5_block3_1_conv (Conv2D)	(None, 2, 2, 512)	1,049,088	conv5_block2_out
conv5_block3_1_bn (BatchNormalization)	(None, 2, 2, 512)	2,048	conv5_block3_1_c
conv5_block3_1_relu (Activation)	(None, 2, 2, 512)	0	conv5_block3_1_t
conv5_block3_2_conv (Conv2D)	(None, 2, 2, 512)	2,359,808	conv5_block3_1_r
conv5_block3_2_bn (BatchNormalization)	(None, 2, 2, 512)	2,048	conv5_block3_2_c
conv5_block3_2_relu (Activation)	(None, 2, 2, 512)	0	conv5_block3_2_t
conv5_block3_3_conv (Conv2D)	(None, 2, 2, 2048)	1,050,624	conv5_block3_2_r

```

1 with open('model_summary.txt', 'w') as f:
2     best_vgg16_model.summary(print_fn=lambda x: f.write(x + '\n'))
3

```



```

1 # Train the best VGG16 model
2 history_vgg16 = best_vgg16_model.fit(
3     datagen.flow(X_train, y_train, batch_size=32),
4     validation_data=(X_val, y_val),
5     epochs=5,
6     callbacks=[early_stopping]
7 )
8
9 # Train the best ResNet50 model
10 history_resnet50 = best_resnet50_model.fit(
11     datagen.flow(X_train, y_train, batch_size=32),
12     validation_data=(X_val, y_val),
13     epochs=5,
14     callbacks=[early_stopping]
15 )
16

```



```

Epoch 1/5
/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:127: UserWarning: `self._warn_if_super_not_called()`
self._warn_if_super_not_called()
119/119 ━━━━━━━━━━━ 33s 183ms/step - accuracy: 0.4654 - loss: 1.9510 - val_accu
Epoch 2/5
119/119 ━━━━━━━━━━━ 18s 51ms/step - accuracy: 0.4613 - loss: 1.8782 - val_accu
Epoch 3/5
119/119 ━━━━━━━━━━━ 8s 66ms/step - accuracy: 0.5049 - loss: 1.7891 - val_accu
Epoch 4/5
119/119 ━━━━━━━━━━━ 9s 52ms/step - accuracy: 0.4990 - loss: 1.7889 - val_accu
Epoch 5/5
119/119 ━━━━━━━━━━━ 8s 65ms/step - accuracy: 0.4981 - loss: 1.7152 - val_accu
Epoch 1/5
119/119 ━━━━━━━━━━━ 35s 172ms/step - accuracy: 0.1882 - loss: 2.5381 - val_accu
Epoch 2/5
119/119 ━━━━━━━━━━━ 8s 66ms/step - accuracy: 0.2030 - loss: 2.4945 - val_accu
Epoch 3/5
119/119 ━━━━━━━━━━━ 10s 62ms/step - accuracy: 0.2178 - loss: 2.4226 - val_accu
Epoch 4/5
119/119 ━━━━━━━━━━━ 9s 53ms/step - accuracy: 0.2065 - loss: 2.4228 - val_accu
Epoch 5/5
119/119 ━━━━━━━━━━━ 10s 53ms/step - accuracy: 0.2145 - loss: 2.3693 - val_accu

```







```

1 # Plot accuracy and loss curves for both models
2 def plot_history(history, model_name):
3     plt.figure(figsize=(12, 4))
4

```

```
5     # Accuracy plot
6     plt.subplot(1, 2, 1)
7     plt.plot(history.history['accuracy'], label='Training Accuracy')
8     plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
9     plt.xlabel('Epoch')
10    plt.ylabel('Accuracy')
11    plt.legend()
12    plt.title(f'{model_name} Model Accuracy')
13
14    # Loss plot
15    plt.subplot(1, 2, 2)
16    plt.plot(history.history['loss'], label='Training Loss')
17    plt.plot(history.history['val_loss'], label='Validation Loss')
18    plt.xlabel('Epoch')
19    plt.ylabel('Loss')
20    plt.legend()
21    plt.title(f'{model_name} Model Loss')
22
23    plt.show()
24
25 plot_history(history_vgg16, "VGG16")
26 plot_history(history_resnet50, "ResNet50")
27
```



```
1 from sklearn.metrics import confusion_matrix, classification_report
2
3 def evaluate_model(model, X_val, y_val, model_name):
4     # Make predictions
5     y_pred = model.predict(X_val)
6     y_pred_classes = np.argmax(y_pred, axis=1)
7     y_true_classes = np.argmax(y_val, axis=1)
8
```

```
9     # Generate confusion matrix
10    conf_matrix = confusion_matrix(y_true_classes, y_pred_classes)
11
12    # Plot confusion matrix
13    plt.figure(figsize=(10, 8))
14    sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
15                xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
16    plt.xlabel('Predicted Labels')
17    plt.ylabel('True Labels')
18    plt.title(f'{model_name} Confusion Matrix')
19    plt.show()
20
21    # Print classification report
22    print(f"{model_name} Classification Report:\n",
23          classification_report(y_true_classes, y_pred_classes, target_names=label_encoder
24                                classes_))
25 # Evaluate both models
26 evaluate_model(best_vgg16_model, X_val, y_val, "VGG16")
27 evaluate_model(best_resnet50_model, X_val, y_val, "ResNet50") # Use the tuned ResNet50 mc
```



30/30

1s 30ms/step




VGG16 Classification Report:

	precision	recall	f1-score	support
Black-grass	0.32	0.38	0.35	47
Charlock	0.43	0.80	0.56	76
Cleavers	0.53	0.68	0.60	62
Common Chickweed	0.65	0.62	0.63	102
Common wheat	0.65	0.34	0.45	38
Fat Hen	0.52	0.45	0.48	94
Loose Silky-bent	0.68	0.74	0.71	141
Maize	0.95	0.43	0.59	42
Scentless Mayweed	0.59	0.47	0.52	105
Shepherds Purse	0.53	0.16	0.25	56
Small-flowered Cranesbill	0.73	0.61	0.66	114
Sugar beet	0.44	0.66	0.53	73
accuracy			0.56	950
macro avg	0.58	0.53	0.53	950
weighted avg	0.59	0.56	0.56	950

30/30

7s 118ms/step



ResNet50 Classification Report:

	precision	recall	f1-score	support
Black-grass	0.00	0.00	0.00	47
Charlock	0.08	0.99	0.15	76
Cleavers	0.06	0.03	0.04	62
Common Chickweed	0.00	0.00	0.00	102
Common wheat	0.25	0.05	0.09	38
Fat Hen	0.00	0.00	0.00	94
Loose Silky-bent	1.00	0.01	0.01	141
Maize	0.00	0.00	0.00	42
Scentless Mayweed	0.00	0.00	0.00	105
Shepherds Purse	0.00	0.00	0.00	56
Small-flowered Cranesbill	0.00	0.00	0.00	114
Sugar beet	0.00	0.00	0.00	73
accuracy			0.08	950
macro avg	0.12	0.09	0.02	950
weighted avg	0.17	0.08	0.02	950

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531: Unde
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```